

Noisy calibration, latent coverage and the location of the mean

BY KUN WOO PARK

[Department, institution, postal address, country]

pkw31386094@gmail.com

SUMMARY

Prediction intervals for a latent quantity, such as the mean of an unsampled area, are often calibrated on noisy proxies with Gaussian errors. The noisy threshold is conservative for symmetric unimodal latent distributions, but symmetry is seldom plausible. For bi-log-concave latent distributions, which include the log-concave ones and some multimodal ones, we show that little is lost without it: at every noise level, noisy coverage 0.9 guarantees latent coverage 0.8991, and 0.8992 cannot be guaranteed, so a slack of one part in a thousand restores the nominal level. Hence the noisy conformal threshold at a suitable rank, at most three above the usual high-probability rank for 29 to 3000 calibration units, covers the latent target with high probability even when the noise variances are heterogeneous and unknown. Behind this is a precise account of the worst-case widening of the calibrated radius, which is governed by the location of the latent mean. With the mean at a fixed distance from the boundary of the interval the widening is of the order of the noise variance, an order attained by asymmetric distributions with mean zero; it reaches the order of the noise standard deviation only when the mean lies within a bounded number of noise standard deviations of the boundary. Under log-concavity, known variances and a numerically computed radius shorten the interval by about ten per cent in simulations.

Some key words: Conformal prediction; Deconvolution; Log-concavity; Measurement error; Small area estimation.

1. INTRODUCTION

Suppose that a residual W , the latent target minus a centre fitted on independent training data, is observed only through $V = W + e$, where $e \sim N(0, D)$ is independent of W and, until § 4.2, D is known. A survey statistician who calibrates a prediction interval for the true mean of an unsampled area on the direct estimates of sampled areas is in this position (Fay & Herriot, 1979), as is a conformal analyst whose calibration responses carry measurement error (Einbinder et al., 2024; Cohen et al., 2025). The threshold t at which the noisy scores $|V|$ have coverage p is observable. We ask which radius gives the latent residual coverage q , uniformly over a class of latent distributions.

If W is symmetric and unimodal, the theorem of Anderson (1955) gives $\text{pr}(|W + e| \leq t) \leq \text{pr}(|W| \leq t)$, so the noisy threshold is conservative (Einbinder et al., 2024). Without symmetry the available transfers make no shape assumption: a median-zero noise turns a noisy level $1 - \alpha$ into a latent level $1 - 2\alpha$ (Guille-Escuret & Ndiaye, 2024), and LatentCP (Zheng et al., 2026), with a marginal guarantee and no shape assumption, uses a Markov-type compatibility threshold. Cohen et al. (2025) estimate the noise-free threshold by deconvolution, without a latent guarantee; conformal intervals for small areas (Bersson & Hoff, 2024) cover observable responses rather than latent area means, and conformal confidence intervals (Zhang & Tuoto, 2026) target the expectations of sampled units rather than a new, unsampled one; Armstrong et al. (2022) and

Ignatiadis & Wager (2022) give robust and nonparametric empirical Bayes intervals for sampled units, and parametric small-area prediction intervals (Hall & Maiti, 2006; Yoshimori & Lahiri, 2014) rely on a model for the latent effects. For a fixed smooth latent distribution, Gaussian noise moves quantiles by $O(D)$ (Jochmans & Weidner, 2024); a guarantee, however, must hold for the worst distribution in a class, and there the order and its cause are different.

Uniformly over a shape class the picture is organized by the location of the mean. We work with log-concave laws and with the larger class of bi-log-concave laws (Dümbgen et al., 2017), which does not require unimodality. Without a restriction on the mean the worst-case widening at equal levels is at most $c_q D^{1/2}$, with a constant that is sharp as $D \rightarrow 0$ (Theorem 1); if the mean lies at distance at least d from the boundary of the interval the widening is at most $D/(2d)$ for every D , an order attained by asymmetric laws with mean zero (Theorem 2 and Supplementary Proposition S1); and a widening of order $D^{1/2}$ requires the mean within $O(D^{1/2})$ of the boundary, with a limiting coefficient given by a one-dimensional extremum (Theorem 3). The widening is also small, at most 1.16% of the radius at $q = 0.9$. Measured in coverage, noisy coverage 0.9 implies latent coverage 0.8991 at every noise level, and 0.8992 cannot be guaranteed (Corollary 1). This is what makes the precise worst case useful: a slack of one part in a thousand absorbs the widening for every noise variance at once, so the noisy threshold at a suitable rank is valid for the latent target when the noise variances are heterogeneous and unknown (Proposition 3), and knowledge of the variances is needed only for efficiency (§ 4.2).

2. THE SHARP RADIUS

Let $\mathcal{L}$ be the class of log-concave laws on the real line, including point masses, and $\mathcal{B} \supset \mathcal{L}$ the class of laws that are point masses or have a continuous distribution function F with $\log F$ and $\log(1 - F)$ concave. Both classes are closed under translation and scaling, and convolution with a log-concave law preserves them (Prékopa, 1973; Saumard, 2019). After scaling t to 1, write $x = D/t^2$ and let $Z \sim N(0, 1)$ be independent of W . Define

$$R_{p,q}(x) = \sup\{Q_q(|W|) : W \in \mathcal{L}, \text{pr}(|W + x^{1/2}Z| \leq 1) \geq p\}, \quad (1)$$

where Q_q is the lower q -quantile, and $R_{p,q}(x) = -\infty$ if nothing is feasible, which happens exactly when $\text{pr}(|x^{1/2}Z| \leq 1) < p$. If calibration certifies noisy coverage at least p at threshold t , then $t R_{p,q}(D/t^2)$ is the smallest radius that guarantees latent coverage q for every log-concave latent law; this is sharpness within the class and for one constraint, not optimality among procedures that use the whole calibration sample. We call $R_{p,q}(x) - 1$ the radius correction. Let $R_{p,q}^{\mathcal{B}}$ be defined by (1) with $\mathcal{B}$ in place of $\mathcal{L}$, so $R_{p,q}^{\mathcal{B}} \geq R_{p,q}$. The upper bounds of §§ 3–4.1 hold for $R_{p,q}^{\mathcal{B}}$; the exact radius $R_{p,q}$ at a fixed x and its computation use log-concavity. For $0 \leq b \leq 1$ let $R_{p,q}^{(b)}(x)$ be defined by (1) with the additional constraint $|E(W)| \leq b$, so that $1 - b$ bounds the distance of the mean from the boundary.

Write $g_x(w) = \text{pr}(|w + x^{1/2}Z| \leq 1)$, so that the constraint in (1) reads $E\{g_x(W)\} \geq p$, and let $\mathcal{F}$ be the family of point masses and of laws with density proportional to $e^{\beta w}$ on a segment $[a, b]$.

PROPOSITION 1. $R_{p,q}(x) = \sup\{Q_q(|W|) : W \in \mathcal{F}, E\{g_x(W)\} \geq p\}$. The same holds for any continuous constraint kernel that vanishes at $\pm\infty$, in particular for the average kernel of § 4.2.

The proof uses the localization theorem of Fradelizi & Guédon (2004) (Supplementary Material); over $\mathcal{F}$ the problem is a two-dimensional maximization of closed forms.

3. WORST-CASE RADIUS CORRECTION AND THE LOCATION OF THE MEAN

80

3.1. A one-sided problem

Near each end of the interval the problem is one-sided. For $u > 0$ let E_u be exponential with rate u , independent of Z , and for the exponential tail $b - E_u$ let

$$m_u(b) = \Pr(b - E_u + Z \leq 0) = \Phi(-b) + e^{-ub+u^2/2} \Phi(b-u),$$

which is continuous and strictly decreasing in b . Let $b_u(p)$ solve $m_u(b) = p$, and write $v_q(u) = b_u(q) - u^{-1} \log(1/q)$ and $k_q(u) = u^{-1} - b_u(q)$. The law $b_u(q) - E_u$ has q -quantile $v_q(u)$ and mean $-k_q(u)$. 85

LEMMA 1. Let $c_{p,q} = \sup\{F_Y^{-1}(q) : Y \in \mathcal{L}, \Pr(Y + Z \leq 0) \geq p\}$ and $c_q = c_{q,q}$. Then $c_{p,q} = \sup_{u>0}\{b_u(p) - u^{-1} \log(1/q)\}$, $c_q > 0$, and $q \mapsto c_q$ is nonincreasing. The same holds with $\mathcal{B}$ in place of $\mathcal{L}$.

LEMMA 2. Let $e^{-1} < q < 1$ and $\kappa \geq 0$. Then 90

$$\sup\{F_Y^{-1}(q) : Y \in \mathcal{B}, E(Y) \leq -\kappa, \Pr(Y + Z \leq 0) \geq q\} = L_q(\kappa) = \sup\{v_q(u) : u > 0, k_q(u) \geq \kappa\}.$$

The function L_q is positive and nonincreasing, and, for $q > 1/2$, $L_q(\kappa) = c_q$ for $\kappa \leq \kappa_q^*$, where κ_q^* is the largest value of k_q on the compact set of maximizers of v_q , and $\kappa L_q(\kappa) \rightarrow \{1 - \log(1/q)\}/2$ as $\kappa \rightarrow \infty$.

Exponential tails are extremal because concavity of $\log F_Y$ limits how fast mass can accumulate below the q -quantile, and they accumulate it as fast as allowed. Both proofs, in the Supplementary 95 Material, replace Y by the exponential tail through a supergradient of $\log F_Y$ at its q -quantile; the replacement lies stochastically below Y , so it keeps the quantile, the noisy constraint and the bound on the mean. Ball arithmetic gives $c_{0.8} = 0.04607$, $c_{0.9} = 0.01906$ and $c_{0.95} = 0.00841$, and in double precision $\kappa_q^* = 6.81, 20.18$ and 50.07 .

3.2. No restriction on the mean 100

THEOREM 1. For every $q \in (0, 1)$, every $x > 0$ and every $W \in \mathcal{B}$ with $\Pr(|W + x^{1/2}Z| \leq 1) \geq q$, $Q_q(|W|) \leq 1 + c_q x^{1/2}$. Moreover $R_{q,q}(x) = 1 + c_q x^{1/2} + o(x^{1/2})$ as $x \rightarrow 0$, and likewise for $R_{q,q}^{\mathcal{B}}$.

Proof of the upper bound. Let $\alpha_R = \Pr(W + x^{1/2}Z > 1)$ and $\alpha_L = \Pr(W + x^{1/2}Z < -1)$, so that $\alpha_L + \alpha_R \leq 1 - q$. The law of $Y = (W - 1)/x^{1/2}$ has $\log F_Y$ concave and $\Pr(Y + Z \leq 0) = 1 - \alpha_R$. By Lemma 1, $F_Y^{-1}(1 - \alpha_R) \leq c_{1-\alpha_R} \leq c_q$, that is, $\Pr(W > 1 + c_q x^{1/2}) \leq \alpha_R$. The same 105 argument applied to $-W$, which uses concavity of $\log(1 - F_W)$, gives $\Pr(W < -1 - c_q x^{1/2}) \leq \alpha_L$, so $\Pr(|W| \leq 1 + c_q x^{1/2}) \geq q$. $\square$

The lower bound, proved in the Supplementary Material, uses $W = 1 + x^{1/2}\{b_u(q + \eta_x) - E_u\}$, whose density rises to a cut-off just outside the threshold and whose mean is within $O(x^{1/2})$ of the boundary. Since $E\{g_x(W)\} \leq g_x(0)$, the constraint is feasible only if $x^{1/2} \leq 1/\Phi^{-1}\{(1+q)/2\}$, 110 so Theorem 1 gives $Q_q(|W|) - 1 \leq c_q/\Phi^{-1}\{(1+q)/2\}$, which is below 1.16% at $q = 0.9$ for every $W \in \mathcal{B}$. For log-concave laws, double-precision upper bounds give $R_{q,q}(x) \leq 1$ for $x \geq 0.013$, hence a correction below 0.22%, and maximization over $\mathcal{F}$ finds 0.063% (Supplementary Material).

3.3. The mean away from the boundary

THEOREM 2. Let $W \in \mathcal{B}$ with $|E(W)| \leq B$, let $\sigma > 0$ and $t > B$, and suppose that $\text{pr}(|W + \sigma Z| \leq t) \geq q$. Then

$$Q_q(|W|) \leq B + \{(t - B)^2 + \sigma^2\}^{1/2} \leq t + \frac{\sigma^2}{2(t - B)}.$$

In particular $R_{p,q}^{(b)}(x) \leq 1 + x/\{2(1 - b)\}$ for all $p \geq q$ and $x > 0$, and $Q_q(|W|) \leq (t^2 + \sigma^2)^{1/2}$ if $E(W) = 0$.

A widening needs a density that rises steeply towards an end of the interval, and a mean far from that end caps the rate of rise. The proof in the Appendix makes this precise with an exponential-tail comparison and the heat equation for quantiles (Steinbrecher & Shaw, 2008), which give $d\{r(s) - B\}^2/ds \geq -1$ for $r(s) = Q_q(|W + s^{1/2}Z|)$. It improves on Theorem 1 only when $x^{1/2} < 2c_q(1 - b)/(1 - c_q^2)$; for mean-zero laws at $q = 0.9$ this means $x < 1.5 \times 10^{-3}$, where both bounds are below 0.08%. The order x cannot be improved, even for mean-zero laws: for laws with density proportional to $e^{\beta w}$ on a segment, centred at their mean, the radius correction is $\frac{1}{2}\beta r_0 \tanh(\beta r_0)x + o(x)$, where $r_0 = Q_q(|W|)$ before scaling, whenever $\pm r_0$ lie inside the support (Supplementary Material). The best coefficient in this family is 0.3466 at $q = 0.8$ and 0.0747 at $q = 0.9$, so the worst case over centred laws lies in $[0.346, 0.5]$ and $[0.0747, 0.5]$. Centring therefore does not replace symmetry.

3.4. The transition

THEOREM 3. Let $1 - e^{-1} < q < 1$, $x > 0$ and $W \in \mathcal{B}$ with $\text{pr}(|W + x^{1/2}Z| \leq 1) \geq q$. Then $|E(W)| < 1$, and if $\Delta = Q_q(|W|) - 1 > 0$,

$$2\{1 - |E(W)|\}\Delta + \Delta^2 \leq x. \quad (2)$$

Moreover, for every $\kappa \geq 0$,

$$\lim_{x \rightarrow 0} x^{-1/2}\{R_{q,q}^{(b_x)}(x) - 1\} = L_q(\kappa), \quad b_x = 1 - \kappa x^{1/2},$$

and the same limit holds with $\mathcal{L}$ replaced by $\mathcal{B}$ in the definition of $R^{(b)}$.

In the original units the mean lies $\kappa D^{1/2}$ from the threshold, and the unrestricted worst case $c_q D^{1/2}$ persists until κ exceeds κ_q^* , about 20 noise standard deviations at $q = 0.9$; beyond it the coefficient decays like $\{1 - \log(1/q)\}/(2\kappa)$ (Fig. 1). Both parts are proved in the Supplementary Material.

4. CALIBRATION SLACK AND A FINITE-SAMPLE PROCEDURE

4.1. Slack removes the correction

The proof of Theorem 1 splits the noisy failure budget between the two ends. With a noisy level $p \geq q$, splitting both the noisy and the latent budgets gives, for every x and every $W \in \mathcal{B}$ feasible at level p , $Q_q(|W|) \leq 1 + C_{p,q}x^{1/2}$, where $C_{p,q} \geq c_{p,q}$, nonincreasing in p , is a worst case over splits of the one-sided constants, defined in the Supplementary Material.

COROLLARY 1. If $C_{p,q} \leq 0$ then $R_{p,q}^{\mathcal{B}}(x) \leq 1$ for every x , and if $c_{p,q} > 0$ then $R_{p,q}(x) > 1$ for all small x . The smallest slack $p - q$ with $C_{p,q} \leq 0$ lies in $(4.3 \times 10^{-3}, 5.0 \times 10^{-3}]$, $(7.9 \times 10^{-4}, 1.0 \times 10^{-3}]$ and $(1.5 \times 10^{-4}, 2.0 \times 10^{-4}]$ at $q = 0.8, 0.9$ and 0.95 , and so does the smallest slack with $R_{p,q}(x) \leq 1$ for all x . In coverage terms, every $W \in \mathcal{B}$ with noisy coverage at least 0.8,




Fig. 1. Leading coefficient $L_q(\kappa)$ of the worst-case radius correction, times 100, when the mean lies at distance at least $\kappa x^{1/2}$ from the boundary (Theorem 3), for $q = 0.8$ (solid), 0.9 (dashes) and 0.95 (dots), with the upper bound $\min\{c_q, (\kappa^2 + 1)^{1/2} - \kappa\}$ from Theorems 1 and 2 (grey). Both axes are logarithmic.

Alt text: Log-log plot of the leading coefficient against the scaled distance of the mean from the boundary for three levels; each curve is flat up to about 7, 20 and 50, then decays like the reciprocal of the distance, just below a grey upper bound.

0.9 or 0.95 has latent coverage at least 0.795, 0.8991 or 0.9498, whatever the noise variance, while 0.7954, 0.8992 and 0.9499 fail for some log-concave laws at small x .

4.2. Heterogeneous and unknown noise variances

Let $W_1, \dots, W_K, W_{\text{new}}$ be independent with law G , independent of $e_i \sim N(0, D_i)$ with the D_i fixed. The scores $V_i = W_i + e_i$ are independent but not identically distributed, and $\mathcal{D}$ denotes the calibration data.

PROPOSITION 2. Let $T = |V|_{(k)}$, let $\bar{F}$ be the average distribution function of the $|V_i|$, assumed continuous, and let p_k be the δ -quantile of the $\text{Beta}(k, K + 1 - k)$ distribution. If $k \geq Kp_k + 1$ then $\text{pr}\{\bar{F}(T) \geq p_k\} \geq 1 - \delta$.

The proof uses Hoeffding (1956); see also Sen (1970).

PROPOSITION 3. Let $G \in \mathcal{B}$, let the D_i be unknown, and suppose that $k \geq Kp_k + 1$ and $C_{p_k, q} \leq 0$. Then $\text{pr}_{\mathcal{D}}\{\text{pr}(|W_{\text{new}}| \leq T \mid \mathcal{D}) \geq q\} \geq 1 - \delta$.

Proof. On the event of Proposition 2, $K^{-1} \sum_i \text{pr}(|W + e_i| \leq T) \geq p_k$ for $W \sim G$. If $\text{pr}(|W| \leq T) < q$, Corollary 1 at threshold T gives $\text{pr}(|W + e_i| \leq T) < p_k$ for every i , whatever D_i , and so for the average, a contradiction. $\square$

At $q = 0.9$ and $\delta = 0.05$ the conditions hold at the smallest rank with $p_k \geq 0.901$ and $k \geq Kp_k + 1$, which exists for $K \geq 29$ and, for $K \leq 3000$, is at most three ranks above the usual high-probability rank, the smallest with $p_k \geq 0.9$ (Supplementary Material). For $K = 110$ the two coincide, at $k = 105$ with $p_k = 0.9068$. The marginal rank $\lceil 0.9(K + 1) \rceil = 100$ is not covered. Known variances shorten the threshold.

COROLLARY 2. Let $G \in \mathcal{B}$ and $k \geq Kp_k + 1$. If $D_i \geq D_{\min}$ for all i and $C_{p', q} < 0$ for some p' with $q \leq p' \leq p_k$, then $s = \max\{0, T + C_{p', q} D_{\min}^{1/2}\}$ satisfies $\text{pr}_{\mathcal{D}}\{\text{pr}(|W_{\text{new}}| \leq s \mid \mathcal{D}) \geq q\} \geq 1 - \delta$, whether or not the D_i are otherwise known.

Proof. On the event of Proposition 2, some i has $\text{pr}(|W + e_i| \leq T) \geq p_k \geq p'$, and the bound of § 4.1 at $x = D_i/T^2$ gives $Q_q(|W|) \leq T + C_{p', q} D_i^{1/2} \leq T + C_{p', q} D_{\min}^{1/2}$. $\square$

Table 1. *Reliability, times 100, and mean width, in parentheses; target reliability 95*

latent law	Markov	Fay–Herriot	noisy threshold	Corollary 2	computed
normal	100 (5.92)	96 (3.93)	100 (4.96)	100 (4.90)	98 (4.45)
Laplace	100 (6.29)	93 (3.96)	100 (5.15)	100 (5.09)	99 (4.66)
centred gamma	100 (6.16)	100 (3.95)	100 (4.97)	100 (4.90)	100 (4.45)
truncated Laplace	100 (5.89)	85 (3.89)	100 (4.93)	100 (4.87)	97 (4.42)
truncated exponential	100 (6.17)	95 (3.93)	100 (5.06)	100 (4.99)	100 (4.55)
bimodal ($\mathcal{B}$ only)	100 (5.77)	100 (4.01)	100 (4.87)	100 (4.81)	99 (4.35)
t_3 (outside $\mathcal{B}$)	100 (5.70)	97 (3.80)	100 (4.66)	100 (4.60)	98 (4.09)

Noisy threshold, T at rank 105; computed, TR^{mix} ; centred gamma with shape 2; truncated exponential, standardized density proportional to e^{4w} on $[0, 1]$; bimodal, the mixture of Supplementary Lemma S1; widths are full widths. The Markov rule, which needs no shape assumption, is guaranteed for every law. Under the bimodal law the noisy threshold and Corollary 2 are guaranteed and the computed radius is exploratory; under t_3 the noisy threshold is guaranteed by Anderson’s theorem, since the law is symmetric and unimodal, while Corollary 2 and the computed radius are not. Standard errors are at most 2.9 for reliabilities and 0.08 for widths.

At $q = 0.9$, $K = 110$ and $\delta = 0.05$, where $p_k = 0.906806$, ball arithmetic gives $C_{0.9068,0.9} \leq -0.114$, hence $s = T - 0.114D_{\min}^{1/2}$; plugging in the smallest variance instead is not valid (Supplementary Material). Among rules that use the single constraint at T , the sharpest use of known variances is the average kernel $\bar{g}(w) = K^{-1} \sum_i \text{pr}(|w + e_i| \leq T)$: $\int \bar{g} dG \geq p_k$ on the event of Proposition 2, Proposition 1 applies for log-concave G , and the computed radius $s = TR_{p_k - \varepsilon, q}^{\text{mix}}(D_1/T^2, \dots, D_K/T^2)$, with ε bounding a discretization error, has the same guarantee.

For the computed radius our implementation returns a numerical upper bound on R^{mix} , by a branch and bound in double precision, not a certificate.

5. NUMERICAL ILLUSTRATION

The target is $\text{pr}_{\mathcal{D}}\{\text{pr}(W_{\text{new}} \in C \mid \mathcal{D}) \geq 0.9\} \geq 0.95$, with $K = 110$, $\text{var}(W) = 1$ and heterogeneous noise variances $D_i = 0.577 L_i / E(L_i)$, L_i lognormal with log-scale 0.7. Latent residuals, all with mean zero, are drawn directly; conditional coverage given each data set is exact, and reliability is estimated from 150 data sets per law. Besides four log-concave laws we use a truncated exponential law in the extremal family $\mathcal{F}$, an asymmetric bimodal mixture in $\mathcal{B}$ but not in $\mathcal{L}$, and a t_3 law, unimodal but outside $\mathcal{B}$. We compare the three rules of § 4.2; a Fay–Herriot tolerance interval, our adaptation of the parametric bootstrap of Chatterjee et al. (2008) (Supplementary Material); and the Markov rule applied to $\bar{g}_T(0) - \bar{g}_T(W) \geq 0$, the smallest radius the single constraint allows without a shape assumption, at its best rank, 107.

The rules with a guarantee reached the target for every law in their class, and the computed radius did so also under the bimodal and t_3 laws, for which it has none. Corollary 2 was 1.2–1.4% shorter than the noisy threshold, and the computed radius 9.5–12% shorter than the noisy threshold and 25–28% shorter than the Markov rule. The Fay–Herriot interval was 7–15% shorter than the computed radius but missed the target under the truncated Laplace law by 5.8 standard errors; its other shortfalls are within Monte Carlo error. With 2000 data sets per law the noisy threshold, Corollary 2 and the Markov rule reached conditional coverage 0.9 in every data set, with mean conditional coverage 0.97–0.99 for the first two, so the guarantees are conservative for laws far from the extremal ones of § 3. At the boundary-layer law that minimizes latent coverage, with noise standard deviation 1% of the threshold and $K = 10^4$, the bounds are tight: the usual high-probability rank had reliability 0.922, standard error 0.003, the certified rank 0.961, and the

marginal rank mean latent coverage 0.89918, between the guaranteed 0.8991 and the unattainable 0.8992 (Supplementary Material). The Supplementary Material compares LatentCP (Zheng et al., 2026), 16–18% longer than the computed radius, and the deconvolution threshold of Cohen et al. (2025), which has no latent guarantee, and shows that at an oracle-matched reliability the gains lie in the cost of a guarantee rather than in the shape of the interval.

6. DISCUSSION

The guarantees are conditional on the calibration sample and assume independent Gaussian noise and a bi-log-concave, or for the computed radius log-concave, latent law shared by the calibration areas and the new one. The target is an unsampled area; for a sampled area the direct estimate is itself informative, and intervals for the expectations of sampled units are a different problem (Zhang & Tuoto, 2026). The shared law is the substantive assumption. In the Supplementary Material, shifting the new residual by half a latent standard deviation lowered mean coverage by up to 0.027 for laws like those of Table 1, but near the extremal law a tenth of a standard deviation breaks the guarantee; and if residual dispersion varies with the noise variance, coverage holds on average over new areas, not given a new area’s variance. Bi-log-concavity cannot be checked on the latent scale, but Gaussian convolution preserves it, so with equal variances a confidence band for the signed noisy residuals containing no bi-log-concave distribution function (Dümbgen et al., 2017) rules it out, though not conversely; it excludes heavy tails, although the symmetric t_3 law of Table 1 is covered by Anderson’s theorem. The mean condition of Theorems 2 and 3 concerns the latent residual law, which an intercept in the fitted regression does not centre. Since $E(V) = E(W)$ the mean is identified, and a bound on $|E(W)|$ with Theorem 2 could correct the radius, but at $q = 0.9$ only where both bounds are below 0.08% of it.

Open questions include non-Gaussian noise, the centred-law coefficient, in $[0.0747, 0.5]$ at $q = 0.9$, a median version of Theorem 2, and whether $C_{p,q} = c_{p,q}$ and $R_{p,q}^{\mathcal{B}} = R_{p,q}$.

DECLARATION OF THE USE OF GENERATIVE AI AND AI-ASSISTED TECHNOLOGIES

[To be completed by the author.]

SUPPLEMENTARY MATERIAL

The Supplementary Material includes omitted proofs, the certificates, the estimated-variance procedure, further simulations and an application. Code reproducing every number is at <https://github.com/kunhailP/noisy-calibration-latent-coverage>.

APPENDIX 1

Proof of Theorem 2

Proof of Theorem 2. Let $\mu = E(W)$. If W is a point mass, $Q_q(|W|) = |\mu| \leq B$. Otherwise, for $s > 0$ the law of $W + s^{1/2}Z$ is in $\mathcal{B}$ (Saumard, 2019), with mean μ , distribution function F_s and a positive smooth density f_s . For $r > \mu$ let $h = f_s(r)/F_s(r)$. The tangent of $\log F_s$ at r gives $F_s(y) \leq \min\{F_s(r)e^{h(y-r)}, 1\}$, the distribution function of an exponential tail with mean $r - \{1 + \log F_s(r)\}/h$ that lies stochastically below $W + s^{1/2}Z$, so $\mu \geq r - \{1 + \log F_s(r)\}/h$ and $h(r - \mu) \leq 1$. Laws in $\mathcal{B}$ satisfy $f'_s \leq f_s^2/F_s$ (Dümbgen et al., 2017), so $(\log f_s)'(r) \leq h \leq 1/(r - \mu)$; applied to $-W$ this gives $(\log f_s)'(-r) \geq -1/(r + \mu)$. Hence, for

$r > B$,

$$f'_s(r) \leq \frac{f_s(r)}{r - B}, \quad -f'_s(-r) \leq \frac{f_s(-r)}{r - B}.$$

245 Let $H_s(r) = \text{pr}(|W + s^{1/2}Z| \leq r)$. By the heat equation $\partial_s f_s = f'_s/2$,

$$\partial_s H_s(r) = \frac{1}{2} \{f'_s(r) - f'_s(-r)\} \leq \frac{f_s(r) + f_s(-r)}{2(r - B)} = \frac{\partial_r H_s(r)}{2(r - B)}.$$

The quantile $r(s) = Q_q(|W + s^{1/2}Z|)$ solves $H_s\{r(s)\} = q$ and is continuously differentiable on $s > 0$, with $r'(s) \geq -1/[2\{r(s) - B\}]$ where $r(s) > B$. Hence $\{r(s) - B\}_+^2$ has derivative at least -1 almost everywhere, and $\{r(\epsilon) - B\}_+^2 \leq \{r(\sigma^2) - B\}_+^2 + \sigma^2 - \epsilon$. As $\epsilon \rightarrow 0$, $r(\epsilon) \rightarrow Q_q(|W|)$, because a nondegenerate law in $\mathcal{B}$ has a positive density on the interior of its support (Dümbgen et al., 2017), so the distribution function of $|W|$ is continuous and strictly increasing where it lies strictly between 0 and 1. Since $r(\sigma^2) \leq t$, the first inequality follows, and the second from $(d^2 + \sigma^2)^{1/2} \leq d + \sigma^2/(2d)$. $\square$

REFERENCES

- ANDERSON, T. W. (1955). The integral of a symmetric unimodal function over a symmetric convex set and some probability inequalities. *Proc. Am. Math. Soc.* **6**, 170–6.
- 255 ARMSTRONG, T. B., KOLESÁR, M. & PLAGBORG-MØLLER, M. (2022). Robust empirical Bayes confidence intervals. *Econometrica* **90**, 2567–602.
- BERSSON, E. & HOFF, P. D. (2024). Optimal conformal prediction for small areas. *J. Surv. Statist. Methodol.* **12**, 1464–88.
- CHATTERJEE, S., LAHIRI, P. & LI, H. (2008). Parametric bootstrap approximation to the distribution of EBLUP and related prediction intervals in linear mixed models. *Ann. Statist.* **36**, 1221–45.
- 260 COHEN, Y., GOLDBERGER, J. & TIRER, T. (2025). Efficient conformal prediction for regression models under label noise. *arXiv: 2509.15120v2*.
- DÜMBGEN, L., KOLESNYK, P. & WILKE, R. A. (2017). Bi-log-concave distribution functions. *J. Statist. Plan. Infer.* **184**, 1–17.
- 265 EINBINDER, B.-S., FELDMAN, S., BATES, S., ANGELOPOULOS, A. N., GENDLER, A. & ROMANO, Y. (2024). Label noise robustness of conformal prediction. *J. Mach. Learn. Res.* **25**, 1–66.
- FAY, R. E. & HERRIOT, R. A. (1979). Estimates of income for small places: an application of James–Stein procedures to census data. *J. Am. Statist. Assoc.* **74**, 269–77.
- FRADELIZI, M. & GUÉDON, O. (2004). The extreme points of subsets of s -concave probabilities and a geometric localization theorem. *Discrete Comput. Geom.* **31**, 327–35.
- 270 GUILLE-ESCURET, C. & NDIAYE, E. (2024). From conformal predictions to confidence regions. *arXiv: 2405.18601*.
- HALL, P. & MAITI, T. (2006). On parametric bootstrap methods for small area prediction. *J. R. Statist. Soc. B* **68**, 221–38.
- HOEFFDING, W. (1956). On the distribution of the number of successes in independent trials. *Ann. Math. Statist.* **27**, 713–21.
- 275 IGNATIADIS, N. & WAGER, S. (2022). Confidence intervals for nonparametric empirical Bayes analysis. *J. Am. Statist. Assoc.* **117**, 1149–66.
- JOCHMANS, K. & WEIDNER, M. (2024). Inference on a distribution from noisy draws. *Economet. Theory* **40**, 60–97.
- PRÉKOPA, A. (1973). On logarithmic concave measures and functions. *Acta Sci. Math. (Szeged)* **34**, 335–43.
- 280 SAUMARD, A. (2019). Bi-log-concavity: some properties and some remarks towards a multi-dimensional extension. *Electron. Commun. Probab.* **24**, Paper no. 61.
- SEN, P. K. (1970). A note on order statistics for heterogeneous distributions. *Ann. Math. Statist.* **41**, 2137–9.
- STEINBRECHER, G. & SHAW, W. T. (2008). Quantile mechanics. *Eur. J. Appl. Math.* **19**, 87–112.
- YOSHIMORI, M. & LAHIRI, P. (2014). A second-order efficient empirical Bayes confidence interval. *Ann. Statist.* **42**, 1233–61.
- 285 ZHANG, L.-C. & TUOTO, T. (2026). Conformal confidence intervals with an application to small area estimation. *arXiv: 2608.02766*.
- ZHENG, M., ZHOU, W. & ZHU, S. (2026). Beyond prediction: conformal inference for latent distributional parameters. *arXiv: 2608.03607*.